ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 1: Why Sierra

Day 1: Why Sierra

Sierra:

- Is the ABS-approved AI assistant
- Uses best-available AI models
- Resides on ABS IT infrastructure
- Protects client data
- Keeps messages private from AI vendors
- Meets compliance requirements

Other AI Assistants:

- Not explicitly approved for use
- Data leakage

ABS
& AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 2: Privacy, Security & Data Governance

Day 2: Privacy, Security & Data Governance

Policy Compliance

- Users must abide by the Sierra Data Policy
- Data Policy is superseded by client contractual terms

App Design & Data Governance

- Chat threads are private
- Workspaces manage collaboration
- Limited web search/external resource usage

Protected IT Infrastructure

- Approved by ABS cyber teams

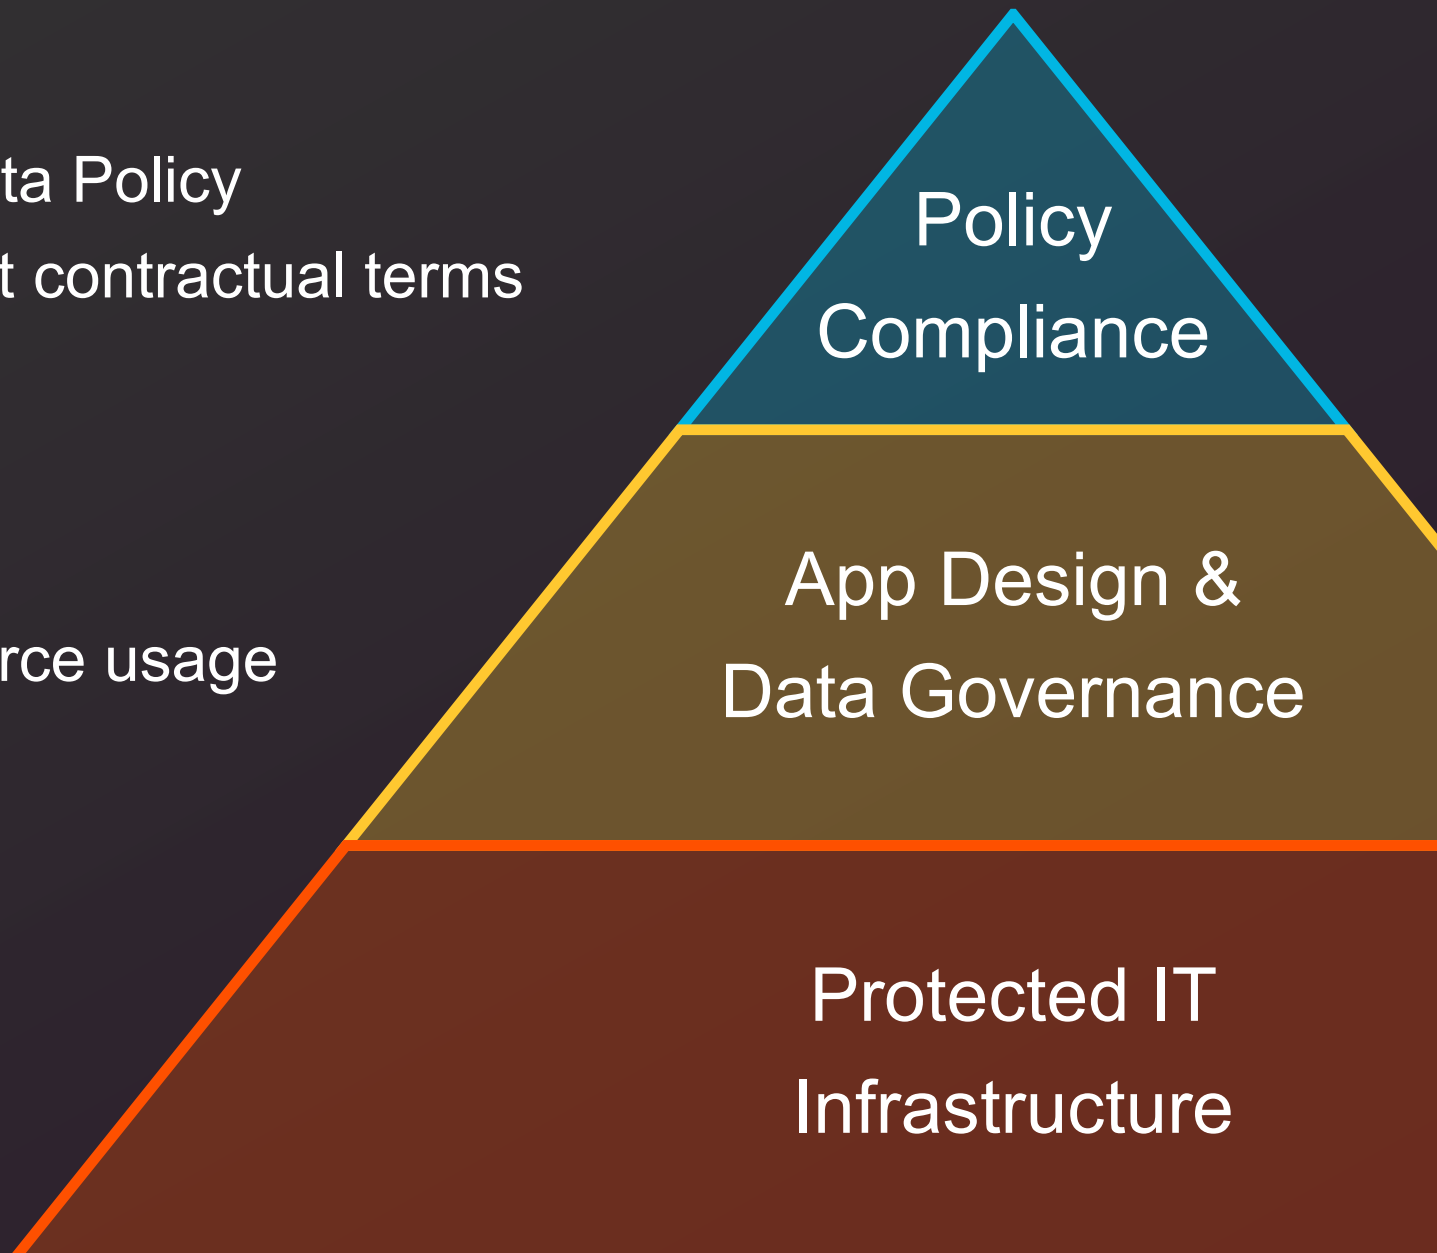

ABS
& AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 3: What is a Large Language Model (LLM)?

Day 3: What is a Large Language Model (LLM)?

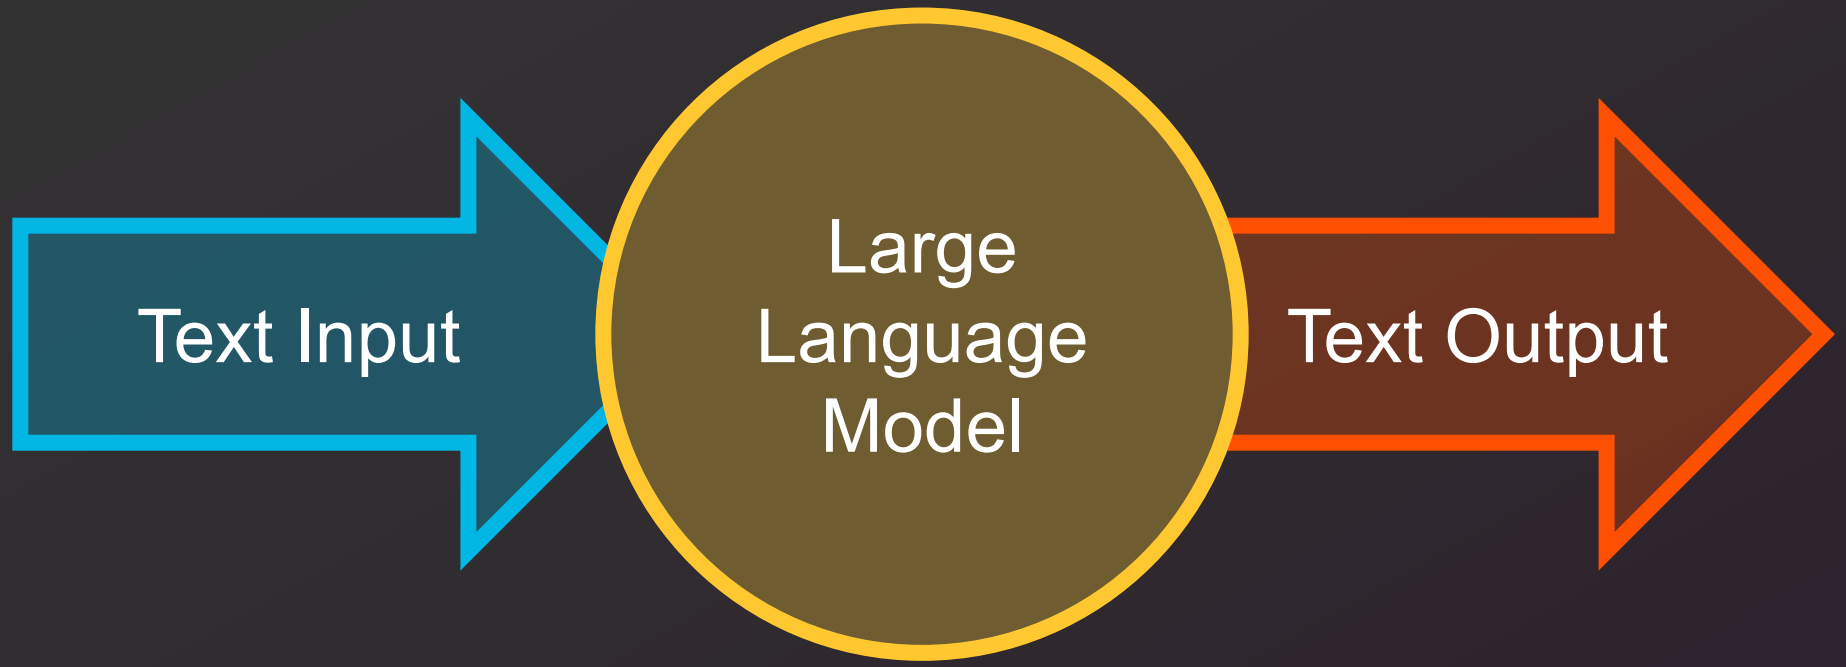

Core Ideas

- An LLM is a calculator that converts a text input (prompt) to a text output (response).
- It has learned patterns from vast amounts of text to produce replies that seem knowledgeable and logical: artificial intelligence.

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 4: What else can LLMs “know”?

Day 4: What else can LLMs “know”?

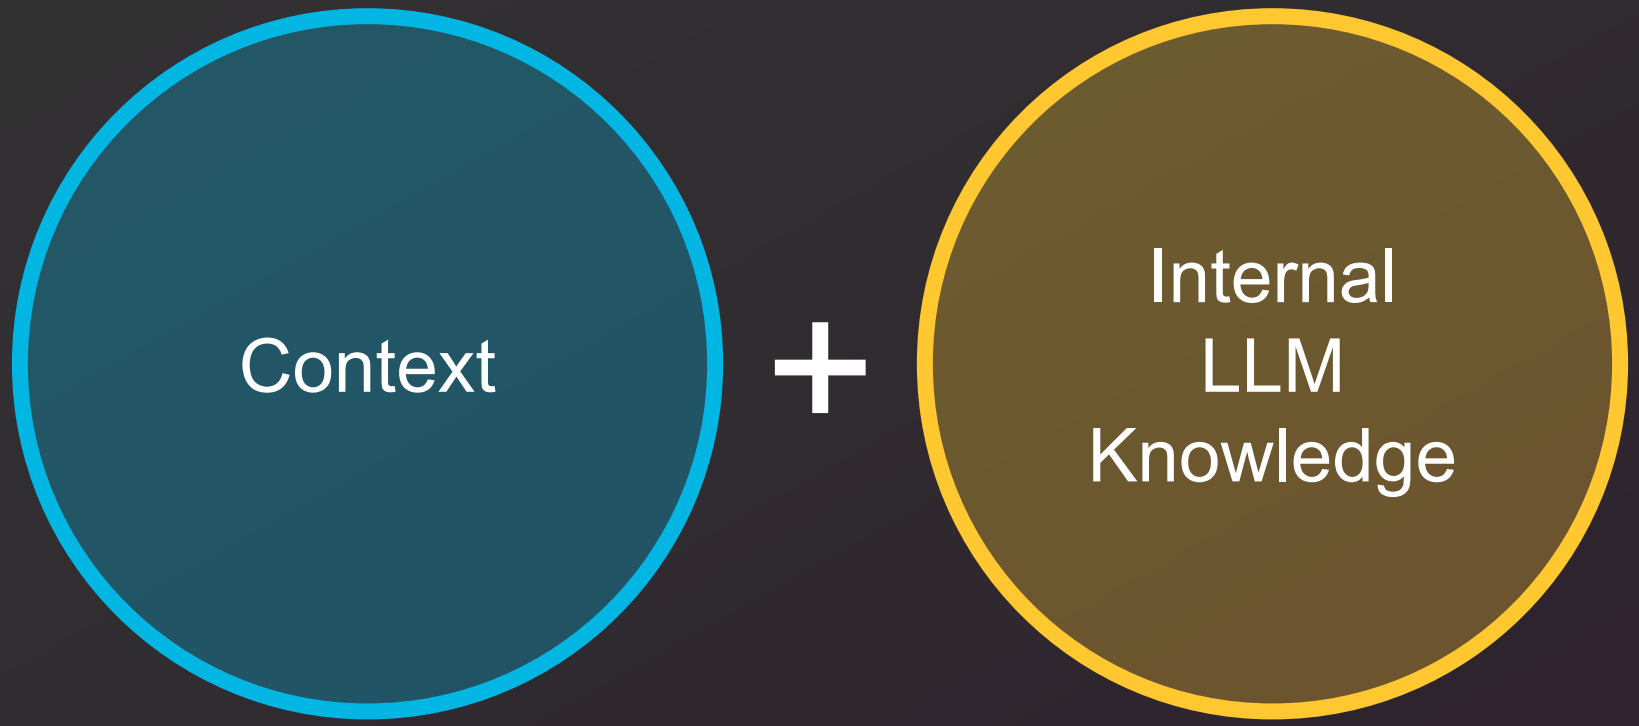

Core Ideas

- The LLM has its own general knowledge that “comes with it.”
- It can also react to external knowledge, “Context,” that is provided to it.
- There is **no** other knowledge available to it

ABS
& AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 5: What LLMs know and do not know

Day 5: What LLMs know and do not know

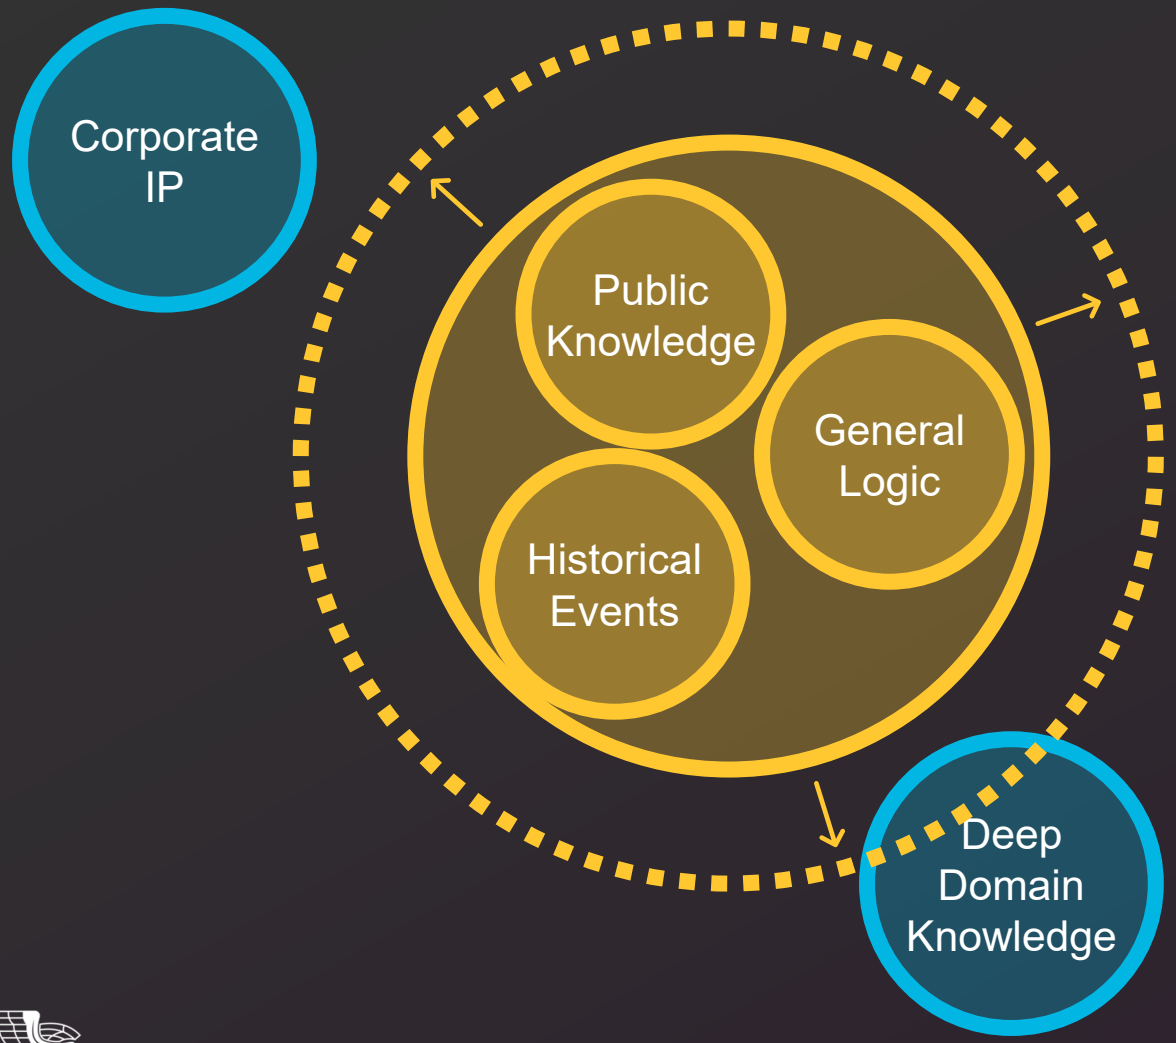

Internal Knowledge Includes:

- Publicly available knowledge (internet sources)
- Events that occurred prior to the creation of their training data (cutoff date)
- Information & logic that is inferable from statistical patterns in the source text
- *Note: newer LLMs have more recent cutoff dates and deeper domain knowledge*

Other knowledge can be provided as context, but is not "built in"

ABS & AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 6: How LLM’s are Built and Customized

Day 6: How LLM's are Built and Customized

Core Ideas

- Training uses massive datasets and compute to create foundation models – this is very expensive!
- Developed by major AI companies: OpenAI, Anthropic, xAI, Google
- LLMs are static – thus we connect them to internal knowledge systems
- ABS does **NOT** train its own LLMs

ABS & AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 7: Tokens & LLM Cost

Day 7: Tokens & LLM Cost

- Every LLM request incurs a cost
- Cost is based on the number of tokens processed – both input and output
- Tokens are small pieces of text – words, parts of words, punctuation
- A token is about 3 -4 characters
- Models process entire conversation history to keep track
- More history = higher cost
- Models perform best when given focused context rather than very long scattered threads

About tokens…

Language models break text into smaller tokens before processing it.

Lan → [1.2323,0.223,-4.797,3.156,…

guage → [-0.156,3.583,0.213,-9.397,…

models → [-4.897,2.196,1.272,…

When you start a fresh topic in Sierra, start a new thread to minimize cost and improve response accuracy.

ABS & AFFILIATED COMPANIES

ABS
& AFFILIATED COMPANIES

SIERRA Literacy Month

Day 8: What is a Hallucination?

Day 8: What is a Hallucination?

Core Ideas

- A “hallucination” is a confident, fluent output not grounded in truth.
- Models are trained to produce the most probable continuation of text, not necessarily the most truthful continuation.
- This often occurs with insufficient context where it’s filling a gap based on patterns but not on real data.
- THIS is why much of our focus is on building internal knowledge systems.

ABS
& AFFILIATED COMPANIES