

SIERRA Literacy Month — Study Guide

Certification Quiz Prep | 22 Days

Day 1 — Why SIERRA

SIERRA is a **secure wrapper** hosting frontier models (ChatGPT, Claude, Grok) in Azure — not its own model, not a fine-tuned GPT. The reason it exists: **data doesn't leave ABS's environment**.

Day 2 — Privacy, Security & Data Governance

Three security layers (bottom to top): **Azure infrastructure** → **app design and data governance** → **policy compliance**. Web search sends only keyword parameters — not your conversation or files. **Client contracts that prohibit AI usage override the ABS data policy**, even on SIERRA.

Day 3 — What Are LLMs?

A model is a **statistical text predictor** — it predicts the most likely continuation of input. It is not retrieving stored facts. Small word changes shift outputs; different models behave differently with the same prompt.

Day 4 — What the Model Knows

Models were trained on public internet text. **ABS internal data was never public — it's not in any model**. Disable external tools and the model can't answer questions about HR policies, recent events, or real-time data.

Day 5 — Context

Context = **your prompt + full conversation history + uploaded files**. The model is static — it reprocesses all available context with each message. Your job is to provide good raw materials.

Day 6 — SIERRA's Tools for Adding Context

Three tools: **file attachments, web search, and ABS knowledge repositories**. The HR tool works because ABS built a knowledge repository — the model wasn't trained on it. **Web search must be explicitly requested** — SIERRA doesn't run it automatically.

Day 7 — Context Windows

One page ≈ **650 tokens**. As threads grow, earlier content gets cut off — the model then falls back on statistical patterns instead of your document. This is a leading cause of hallucinations. **Start new threads when topics change**.

Day 8 — Hallucinations

A hallucination is **confident, fluent output that is factually wrong**. Polished writing is not a signal of accuracy. Verify before using SIERRA output in client deliverables.

Day 9 — Prompting Fundamentals

Four components: **Role, Task, Context, Format**. Prompting is iterative — your first message is a starting point. **Highest-impact approach: ask SIERRA to write the prompt for you** from a rough description of what you need.

Day 10 — Structuring Complex Tasks

The most common mistake: asking for the final deliverable without completing upstream steps. Break the task into steps, review each output before proceeding. After a successful workflow, ask SIERRA to **produce a reusable prompt from the full conversation**.

Day 11 — Structuring Prompts for Data Extraction

Define before you write: **Field → Definition → Exclusions → Location → Decision rules → Evidence → Output format**. Evidence = ask the model to return a quote or clause reference alongside each extracted value. **Phased extraction**: Pass 1 pulls high-level anchors; Pass 2 injects those anchors and goes granular.

Day 12 — Working with Files, Part 1

File flow: **upload → convert → index into vector store → query-based retrieval into context**. Different queries surface different sections. **Files don't persist across threads** — re-upload in each new conversation.

Day 13 — Working with Files, Part 2

First thing to check when output is wrong: preview the conversion, not the prompt. Fast = text extraction (default). Full = OCR for scanned/image-heavy files. The inputs shape the outputs.

Day 14 — Memories

Memories are **persistent context injected into every chat automatically** — set once, always available. Best use: stable information about you (role, specialization, region, output preferences). Not for project-specific or time-sensitive content.

Day 15 — Workspaces

Thread content = private, temporary. **Workspace content = shared and persistent** for all members: context, memories, files, task modes. A file uploaded to a workspace is available to every member without re-upload.

Day 16 — Workspace Task Modes

Context = always-on background for all threads. **Task modes** = invoked for specific, repeatable jobs. Rule: background info → context; step-by-step instructions → task mode. If you've written the same multi-step prompt twice, it belongs in a task mode.

Day 17 — Workspace Walkthrough

Five tabs: **Settings, Context, Files, Task Modes, Memories**. **Personal memories** = about you (role, preferences) — yours alone. **Workspace memories** = about the project or client — shared with all members.

Day 18 — Building a Workspace

Start in a **regular chat**, not the workspace. Use SIERRA to analyze sample materials and draft the context and task mode instructions — then move them in. The models have no idea what a SIERRA workspace is; you have to explain the structure. Design task modes to **pause and confirm key variables** before generating output — don't let it run end-to-end without a checkpoint.

Day 19 — Plato Demo

Plato synthesizes across **multiple rule sections** to answer scenario-based, multi-part questions — ask it the way you'd ask a colleague. Follow-up questions that add scenario context produce more targeted answers. Always verify citations.

Day 20 — Environmental Plan Review Workspace

One workspace, multiple task modes — shared context lives at the workspace level; specific review instructions live in individual task modes. Build task mode instructions by pasting your criteria into a regular chat and asking SIERRA to turn them into a task mode. Good review output includes: **yes/no flags, findings, source references, and a reviewer check column**.

Day 21 — Web Search & Deep Research

Web search: fast, structured summary with citations. **Deep research**: comprehensive multi-source report, downloadable. Both require source verification — click through before using the output.

Day 22 — HR Policy Repository

SIERRA reads your **entity and home country** from your profile and applies the correct policy automatically — you don't specify it. Two employees in different countries get different answers to the same question. If a policy isn't in the repository, SIERRA says so and directs you to HR — it doesn't guess.

Things People Get Wrong

- **"SIERRA knows my file"** — It indexed it. Different questions surface different sections. Files don't persist across threads.
- **"Well-written = accurate"** — No. Verify before using in deliverables.
- **"One long thread is fine"** — Long threads degrade quality. Start fresh when topics change.
- **"I need to specify my country for HR questions"** — SIERRA reads your profile automatically.
- **"I need to write workspace components from scratch"** — Use a regular chat to draft them.
- **"Web search runs automatically"** — You have to explicitly request it.
- **"Context and task modes are the same"** — Context is always on; task modes are invoked for specific jobs.